

# Designing a 125 GB Monolithic Software Package — First 125 MB as Executable; Remainder as Assets/Data — Target Environments: Windows 10/11 (native) and Browser-only WebAssembly

## 1. Concept Overview

This document proposes a packaging convention in which a single 125 GB file contains two logical parts:

- The first 125.00 MB is the executable (the OS-loadable program and its core runtime code).
- Everything after 125 MB consists of assets (models, media, maps), databases, and configuration payloads.

On Windows 10/11, this can be implemented using the Portable Executable (PE) format with an appended overlay region.

In the browser, a WebAssembly (WASM) core (~125 MB) boots the app, which then streams and manages the remaining assets under a virtual filesystem backed by the web platform (Cache Storage/OPFS/IndexedDB).

## 2. Windows 10/11 Native Implementation

### 2.1. PE/COFF and the “Overlay”

- Windows loads only the sections declared in the PE headers. Any bytes after the last section—often called the overlay—are ignored by the OS loader.
- Therefore, a single file can contain a normal signed PE at the start, followed by arbitrary data. Installers and self-extracting archives already use this pattern at smaller scales.

### 2.2. Enforcing the 125 MB Rule

- Place the entire PE (headers + .text/.rdata/.data/.rsrc/.reloc, etc.) within the first 125 MB. The rest of the file is the overlay.
- Inside the executable, embed a tiny, early-load manifest (TOC) that describes where assets live in the overlay (offsets, sizes, hashes, MIME, target install path, compression, encryption).
- Alignment: PE SectionAlignment and FileAlignment govern mapped code/data; the overlay does not need PE alignment but should keep large-block alignment (e.g., 4 KiB or 64 KiB) to improve paging and I/O.

### 2.3. Startup Flow

- 1) OS maps the 125 MB PE into memory; overlay remains on disk.
- 2) Your bootstrap code parses the embedded TOC and validates the package signature.
- 3) It either (a) memory-maps slices of the overlay on-demand, or (b) extracts assets into proper locations, e.g.:
  - Immutable binaries: %ProgramFiles%\Vendor\App\bin
  - Large content: %ProgramData%\Vendor\App\assets (shared) or %LOCALAPPDATA%\Vendor\App\assets (per-user)
  - Mutable DBs: %LOCALAPPDATA%\Vendor\App\data

4) Verify integrity (SHA-256 per chunk), decompress (zstd/LZMA), and decrypt (if used) before first use.

5) Start the main app with a deterministic search path (private “DLL directory”, SxS assemblies) to avoid DLL hijacking.

#### 2.4. DLLs and Side-by-Side (SxS)

- Prefer statically linking critical runtime or package private DLLs placed next to the EXE (or in an app-local SxS). Use an app manifest to pin versions and enable SafeDllSearchMode.
- Avoid placing DLLs into system directories. Keep everything app-local to ease updates and reduce hijacking risk.

#### 2.5. Code Signing, SmartScreen, and AV

- Authenticode covers the entire file (excluding the signature blob). Appending an overlay must be done *before* signing; any post-signing modification invalidates the signature.
- Submit to Microsoft for reputation building to avoid SmartScreen prompts. Expect long scans for a 125 GB file; prefer delta updates to keep downloads smaller.

#### 2.6. Updating a 125 GB Package

- Use block-level deltas (Rsync/Rolling hash; bsdiff/xdelta; or a custom chunker like 4–16 MiB content-defined chunks).
- Maintain a content-addressed store on disk so unchanged chunks are reused across versions.
- Use BITS or a resilient updater service to handle resume and throttling.

#### 2.7. Filesystem, Limits, and Performance

- NTFS supports 125 GB files; avoid FAT32 media. Enable LargePages only for mapped sections that benefit.
- Memory footprint remains modest because the OS maps only PE sections; overlay is streamed.
- Consider per-title caching and prefetching hints (SuperFetch is deprecated, but your bootstrap can model usage to prefetch critical ranges).

## 3. Browser-Only WebAssembly Implementation

#### 3.1. What “Executable” Means on the Web

- There are no Windows DLLs in a browser. Your 125 MB ‘executable’ is a WASM core module (possibly split into multiple modules) plus a minimal JS bootstrap.
- Use streaming compilation/instantiation so the runtime begins before the full 125 MB arrives.

#### 3.2. Module Layout and Linking

- Split into components: core.wasm (~125 MB) + feature modules (graphics.wasm, audio.wasm, sim.wasm). Use the WebAssembly Component Model or dynamic linking conventions supported by your toolchain.
- Threads: use WASM threads via SharedArrayBuffer; ensure COOP/COEP headers and cross-origin isolation.
- Memory: today’s WASM32 limits single-module linear memory to ~4 GiB; WASM64 (memory64) lifts this ceiling as browsers adopt it. Architect for tiling/streaming so you never require multi-GiB contiguous buffers.

#### 3.3. Delivering the Remaining ~125 GB of Assets

- Do not ship a literal 125 GB single HTTP object. Instead:
  - Chunk assets into 4–16 MiB (CDN-friendly). Name by content hash (e.g., SHA-256) for dedup

and integrity.

- Use Range requests and parallel fetch for large objects.
- Compress with zstd/brotli; serve over HTTP/3 with proper prioritization.
- Storage:
  - Cache Storage for immutable chunks.
  - OPFS (Origin Private File System) for DBs and write-heavy content.
  - IndexedDB for metadata/TOC if convenient.
- Include a signed manifest (TOC) with per-chunk size/hash and a top-level signature (Detached CMS or COSE). Verify in the WASM core before accepting data.

### 3.4. Database Options

- SQLite-WASM or DuckDB-WASM with OPFS backends provide robust embedded DBs. Keep write-amplification in mind; prefer append-only or columnar formats for huge analytical data.

### 3.5. Packaging and Install Experience

- Make it a PWA: installable, offline-capable. A Service Worker manages caching, updates, and background sync (when the app is open).
- First-run wizard lets users choose storage quota (where supported), data packs (language, regions), and download schedules.

### 3.6. Security

- Use HTTPS + HSTS. Set strict CSP to block unexpected script execution. Use Subresource Integrity (SRI) for static JS; for WASM and chunked assets, verify hashes and a signed manifest inside the app.
- Cross-origin isolation (COOP/COEP) for threads and high-resolution timers.

### 3.7. Updates in the Browser

- The Service Worker fetches a new manifest; only changed chunks download. Maintain rolling caches; evict least-recently-used packs.
- Offer fast-path micro-updates for the WASM core (e.g., separate small patch modules) to avoid re-downloading the full 125 MB core.

### 3.8. What You Cannot Do in Browser-Only Mode

- No direct use of Windows DLLs or registry. No arbitrary file system access outside the origin sandbox. No kernel drivers. Integrations must use web APIs (WebGPU/WebAudio/WebUSB where available).

## 4. Shared Container Format (Native & Web)

Define a single content format so both the Windows loader and the WASM runtime read the same manifests and chunks.

Header (within first 125 MB):

- Magic: QPF0; Version; Flags (endianness, compression, encryption)
- TOC offset/length (points into the overlay)
- Public key ID(s) and signature block offset for whole-file or per-chunk signing
- Bootstrapping table (names of primary executable/WASM module(s), entrypoints)

Overlay (after 125 MB):

- Chunks: [hash | offset | length | compression | encryption | targetPath | MIME]
- Optional pack files (e.g., .pak) with internal LZ4/zstd and their own mini-TOC

- End of file trailer repeating critical metadata for recovery

Operational notes:

- Place a small TOC cache in the first 125 MB for cold start discovery; the full TOC can live in the overlay.
- Keep chunk sizes CDN friendly; align to 1–4 MiB boundaries to optimize HTTP/3 and disk I/O.

## 5. Practical Constraints and Recommendations

- Distribution: 125 GB is immense. Prefer optional content packs and regional/language downloads. Use a CDN with partial object and range request optimization.
- Telemetry: Log chunk cache hits/misses (privacy preserving) to tune prefetch and storage budgets.
- Legal/Compliance: Shipping cryptography may trigger export rules; include a crypto notice if using encryption.
- Testing: Build a synthetic content generator to fuzz TOC and chunk boundaries; test power loss recovery and partial downloads.
- Accessibility: Ensure the bootstrap supports screen readers and keyboard navigation even before full assets are present.

## 6. Sample TOC Entry (JSON)

```
{
  "version": 3,
  "chunks": [
    {
      "path": "assets/terrain/region-17.pack",
      "sha256": "",
      "size": 16777216,
      "offset": 134217728,
      "compression": "zstd",
      "encryption": "none",
      "mime": "application/octet-stream"
    },
    {
      "path": "bin/plugins/physics.wasm",
      "sha256": "",
      "size": 24431872,
      "offset": 987654321,
      "compression": "none",
      "encryption": "none",
      "mime": "application/wasm"
    }
  ],
  "signatures": {
    "algorithm": "ed25519",
    "manifestSig": ""
  }
}
```

## 7. Implementation Checklist

- Keep the PE within 125 MB; everything else is overlay.
- Embed a minimal TOC pointer in the executable; store the full TOC in the overlay.
- Sign after the overlay is appended; verify signatures at startup.
- Use app local DLLs (or static link); lock DLL search path via manifest.
- In the browser, ship a ~125 MB core WASM with streaming instantiation and split feature modules.
- Store assets via Service Worker in Cache Storage/OPFS; verify per chunk hashes.
- Use chunked, content addressed assets with delta updates.
- Provide a first run downloader with quotas, selectable packs, and resume.
- Test crashes/power loss during extraction and during web cache population.
- Monitor performance; keep hot assets near the start of the overlay or in separate packs.

## Appendix A: Registry and Configuration on Windows

- Prefer per-user configuration under HKCU\Software\Vendor\App and files under %LOCALAPPDATA%.
- Avoid HKLM writes unless the installer is elevated (MSIX or MSI). If using MSIX, you can keep the one-file packaging for download, but install into MSIX layout at rest.
- Use Event Tracing for Windows (ETW) channels for diagnostics without excessive disk logging.

## Appendix B: Example Startup Pseudocode

```
if !verify_signature(exe_and_overlay): fail("Signature invalid")
toc = read_toc(pointer_in_exe)
for chunk in required_boot_chunks(toc):
    map_or_extract(chunk)
launch_main()
background_prefetch(usage_model, network_budget)
```